

# Hamming weight

The **Hamming weight** of a [string](#) is the number of symbols that are different from the zero-symbol of the [alphabet](#) used. It is thus equivalent to the [Hamming distance](#) from the all-zero string of the same length. For the most typical case, a string of [bits](#), this is the number of 1's in the string. In this binary case, it is also called the **population count**, **popcount**, or **sideways sum**.<sup>[1]</sup> It is the [digit sum](#) of the [binary representation](#) of a given number and the [ℓ<sub>1</sub> norm](#) of a bit vector.

## History and usage

The Hamming weight is named after [Richard Hamming](#) although he did not originate the notion.<sup>[2]</sup> Irving S. Reed introduced a concept, equivalent to Hamming weight in the binary case, in 1954.<sup>[3]</sup>

Examples	
string	Hamming weight
11101	4
11101000	4
00000000	0
789012340567	10

Hamming weight is used in several disciplines including [information theory](#), [coding theory](#), and [cryptography](#). Examples of applications of the Hamming weight include:

- In modular [exponentiation by squaring](#), the number of modular multiplications required for an exponent  $e$  is  $\log_2 e + \text{weight}(e)$ . This is the reason that the public key value  $e$  used in [RSA](#) is typically chosen to be a number of low Hamming weight.
- The Hamming weight determines path lengths between nodes in [Chord distributed hash tables](#).<sup>[4]</sup>
- [IrisCode](#) lookups in biometric databases are typically implemented by calculating the [Hamming distance](#) to each stored record.
- In [computer chess](#) programs using a [bitboard](#) representation, the Hamming weight of a bitboard gives the number of pieces of a given type remaining in the game, or the number of squares of the board controlled by one player's pieces, and is therefore an important contributing term to the value of a position.
- Hamming weight can be used to efficiently compute [find first set](#) using the identity  $\text{ffs}(x) = \text{pop}(x \wedge (\sim(-x)))$ . This is useful on platforms such as [SPARC](#) that have hardware Hamming weight instructions but no hardware find first set instruction.<sup>[5]</sup>
- The Hamming weight operation can be interpreted as a conversion from the [unary numeral system](#) to [binary numbers](#).<sup>[6]</sup>
- In implementation of some [succinct data structures](#) like [bit vectors](#) and [wavelet](#)

[trees](#).

## Efficient implementation

The population count of a [bitstring](#) is often needed in cryptography and other applications. The [Hamming distance](#) of two words  $A$  and  $B$  can be calculated as the Hamming weight of  $A \oplus B$ .

The problem of how to implement it efficiently has been widely studied. Some processors have a single command to calculate it (see below), and some have parallel operations on bit vectors. For processors lacking those features, the best solutions known are based on adding counts in a tree pattern. For example, to count the number of 1 bits in the 16-bit binary number  $a = 0110\ 1100\ 1011\ 1010$ , these operations can be done:

Expression	Binary	Decimal	Comment
$a$	01 10 11 00 10 11 10 10		The original number
$b_0 = (a \gg 0) \& 01\ 01\ 01\ 01$ $01\ 01\ 01\ 01$	01 00 01 00 00 01 00 00	1,0,1,0,0,1,0,0	every other bit from $a$
$b_1 = (a \gg 1) \& 01\ 01\ 01\ 01$ $01\ 01\ 01\ 01$	00 01 01 00 01 01 01 01	0,1,1,0,1,1,1,1	the remaining bits from $a$
$c = b_0 + b_1$	01 01 10 00 01 10 01 01	1,1,2,0,1,2,1,1	list giving # of 1s in each 2-bit slice of $a$
$d_0 = (c \gg 0) \& 0011\ 0011$ $0011\ 0011$	0001 0000 0010 0001	1,0,2,1	every other count from $c$
$d_2 = (c \gg 2) \& 0011\ 0011$ $0011\ 0011$	0001 0010 0001 0001	1,2,1,1	the remaining counts from $c$
$e = d_0 + d_2$	0010 0010 0011 0010	2,2,3,2	list giving # of 1s in each 4-bit slice of $a$
$f_0 = (e \gg 0) \& 00001111$ $00001111$	00000010 00000010	2,2	every other count from $e$
$f_4 = (e \gg 4) \& 00001111$ $00001111$	00000010 00000011	2,3	the remaining counts from $e$
$g = f_0 + f_4$	00000100 00000101	4,5	list giving # of 1s in each 8-bit slice of $a$
$h_0 = (g \gg 0) \& 0000000011111111$	0000000000000101	5	every other count from $g$
$h_8 = (g \gg 8) \& 0000000011111111$	0000000000000100	4	the remaining counts from $g$
$i = h_0 + h_8$	0000000000001001	9	the final answer of the

		16-bit word
--	--	-------------

Here, the operations are as in [C](#), so  $x \gg y$  means to shift  $x$  right by  $y$  bits,  $x \& y$  means the bitwise AND of  $x$  and  $y$ , and  $+$  is ordinary addition. The best algorithms known for this problem are based on the concept illustrated above and are given here:

*//types and constants used in the functions below*

```
const uint64_t m1  = 0x5555555555555555; //binary: 0101...
const uint64_t m2  = 0x3333333333333333; //binary: 00110011..
const uint64_t m4  = 0x0f0f0f0f0f0f0f0f; //binary:  4 zeros,  4 ones ...
const uint64_t m8  = 0x00ff00ff00ff00ff; //binary:  8 zeros,  8 ones ...
const uint64_t m16 = 0x0000ffff0000ffff; //binary: 16 zeros, 16 ones ...
const uint64_t m32 = 0x00000000ffffffff; //binary: 32 zeros, 32 ones
const uint64_t hff = 0xffffffffffffffff; //binary: all ones
const uint64_t h01 = 0x0101010101010101; //the sum of 256 to the power of 0,1,2,3...
```

*//This is a naive implementation, shown for comparison,*

*//and to help in understanding the better functions.*

*//It uses 24 arithmetic operations (shift, add, and).*

```
int popcount_1(uint64_t x) {
    x = (x & m1 ) + ((x >> 1) & m1 ); //put count of each 2 bits into those 2 bits
    x = (x & m2 ) + ((x >> 2) & m2 ); //put count of each 4 bits into those 4 bits
    x = (x & m4 ) + ((x >> 4) & m4 ); //put count of each 8 bits into those 8 bits
    x = (x & m8 ) + ((x >> 8) & m8 ); //put count of each 16 bits into those 16 bits
    x = (x & m16) + ((x >> 16) & m16); //put count of each 32 bits into those 32 bits
    x = (x & m32) + ((x >> 32) & m32); //put count of each 64 bits into those 64 bits
    return x;
}
```

*//This uses fewer arithmetic operations than any other known*

*//implementation on machines with slow multiplication.*

*//It uses 17 arithmetic operations.*

```
int popcount_2(uint64_t x) {
    x -= (x >> 1) & m1; //put count of each 2 bits into those 2 bits
    x = (x & m2) + ((x >> 2) & m2); //put count of each 4 bits into those 4 bits
    x = (x + (x >> 4)) & m4; //put count of each 8 bits into those 8 bits
    x += x >> 8; //put count of each 16 bits into their lowest 8 bits
    x += x >> 16; //put count of each 32 bits into their lowest 8 bits
    x += x >> 32; //put count of each 64 bits into their lowest 8 bits
    return x & 0x7f;
}
```

*//This uses fewer arithmetic operations than any other known*

*//implementation on machines with fast multiplication.*

```
//It uses 12 arithmetic operations, one of which is a multiply.
int popcount_3(uint64_t x) {
    x -= (x >> 1) & m1;           //put count of each 2 bits into those 2 bits
    x = (x & m2) + ((x >> 2) & m2); //put count of each 4 bits into those 4 bits
    x = (x + (x >> 4)) & m4;       //put count of each 8 bits into those 8 bits
    return (x * h01)>>56; //returns left 8 bits of x + (x<<8) + (x<<16) + (x<<24) + ...
}
```

The above implementations have the best worst-case behavior of any known algorithm. However, when a value is expected to have few nonzero bits, it may instead be more efficient to use algorithms that count these bits one at a time. As Wegner (1960) described,<sup>[7]</sup> the [bitwise and](#) of  $x$  with  $x - 1$  differs from  $x$  only in zeroing out the least significant nonzero bit: subtracting 1 changes the rightmost string of 0s to 1s, and changes the rightmost 1 to a 0. If  $x$  originally had  $n$  bits that were 1, then after only  $n$  iterations of this operation,  $x$  will be reduced to zero. The following implementation is based on this principle.

```
//This is better when most bits in x are 0
//It uses 3 arithmetic operations and one comparison/branch per "1" bit in x.
int popcount_4(uint64_t x) {
    int count;
    for (count=0; x; count++)
        x &= x-1;
    return count;
}
```

If we are allowed greater memory usage, we can calculate the Hamming weight faster than the above methods. With unlimited memory, we could simply create a large lookup table of the Hamming weight of every 64 bit integer. If we can store a lookup table of the hamming function of every 16 bit integer, we can do the following to compute the Hamming weight of every 32 bit integer.

```
static uint8_t wordbits[65536] = { /* bitcounts of integers 0 through 65535, inclusive */ }
static int popcount(uint32_t i)
{
    return (wordbits[i&0xFFFF] + wordbits[i>>16]);
}
```

## Language support

Some C compilers provide intrinsics that provide bit counting facilities. For example, [GCC](#) (since version 3.4 in April 2004) includes a builtin function `__builtin_popcount` that will use a processor instruction if available or an efficient library implementation otherwise.<sup>[8]</sup> [LLVM-GCC](#) has included this function since version 1.5 in June, 2005.<sup>[9]</sup>

In [C++ STL](#), the bit-array data structure `bitset` has a `count()` method that counts the number of bits that are set.

In Java, the growable bit-array data structure [BitSet](#) has a [BitSet.cardinality\(\)](#) method that counts the number of bits that are set. In addition, there are [Integer.bitCount\(int\)](#) and [Long.bitCount\(long\)](#) functions to count bits in primitive 32-bit and 64-bit integers, respectively. Also, the [BigInteger](#) arbitrary-precision integer class also has a [BigInteger.bitCount\(\)](#) method that counts bits.

In [Common Lisp](#), the function `logcount`, given a non-negative integer, returns the number of 1 bits. (For negative integers it returns the number of 0 bits in 2's complement notation.) In either case the integer can be a `BIGNUM`.

Starting in [GHC 7.4](#), the [Haskell](#) base package has a `popCount` function available on all types that are instances of the `Bits` class (available from the `Data.Bits` module).<sup>[10]</sup>

[MySQL](#) version of [SQL](#) language provides `BIT_COUNT()` as a standard function.<sup>[11]</sup>

[Fortran 2008](#) has the standard, intrinsic, elemental function `popcnt` returning the number of nonzero bits within an integer (or integer array), see page 380 in *Metcalfe, Michael; John Reid; Malcolm Cohen (2011). Modern Fortran Explained. Oxford University Press. ISBN 0-19-960142-9.*

## Processor support

- [Cray](#) supercomputers early on featured a population count [machine instruction](#), rumoured to have been specifically requested by the U.S. government [National Security Agency](#) for [cryptanalysis](#) applications.
- [AMD's Barcelona](#) architecture introduced the abm (advanced bit manipulation) [ISA](#) introducing the `POPCNT` instruction as part of the [SSE4a](#) extensions.

- [Intel Core](#) processors introduced a POPCNT instruction with the [SSE4.2 instruction set](#) extension, first available in a [Nehalem](#)-based [Core i7](#) processor, released in November 2008.
- [Compaq's Alpha 21264A](#), released in 1999, was the first Alpha series CPU design that had the count extension (CIX).
- [Donald Knuth](#)'s model computer [MMIX](#) that is going to replace [MIX](#) in his book [The Art of Computer Programming](#) has an `SADD` instruction. `SADD a,b,c` counts all bits that are 1 in `b` and 0 in `c` and writes the result to `a`.
- The [ARM architecture](#) introduced the VCNT instruction as part of the Advanced SIMD (NEON) extensions.
- [Analog Devices' Blackfin](#) processors feature the ONES instruction to perform a 32-bit population count.

## See also

- [Minimum weight](#)
- [Two's complement](#)
- [Most frequent k characters](#)

## References

1. ^ *D. E. Knuth (2009). The Art of Computer Programming Volume 4, Fascicle 1: Bitwise tricks & techniques; Binary Decision Diagrams. Addison–Wesley Professional. ISBN 0-321-58050-8.* Draft of [Fascicle 1b](#) available for download.
2. ^ *Thompson, Thomas M. (1983), From Error-Correcting Codes through Sphere Packings to Simple Groups, The Carus Mathematical Monographs #21, The Mathematical Association of America, p. 33*
3. ^ *Reed, I.S. (1954), "A Class of Multiple-Error-Correcting Codes and the Decoding Scheme", I.R.E. (I.E.E.E.), PGIT-4: 38*
4. ^ *Stoica, I., Morris, R., Liben-Nowell, D., Karger, D. R., Kaashoek, M. F., Dabek, F., and Balakrishnan, H. Chord: a scalable peer-to-peer lookup protocol for internet applications. IEEE/ACM Trans. Netw. 11, 1 (Feb. 2003), 17-32. Section 6.3: "In general, the number of fingers we need to follow will be the number of ones in the binary representation of the distance from node to query."*
5. ^ *SPARC International, Inc. (1992). The SPARC architecture manual : version 8 (PDF) (Version 8. ed.). Englewood Cliffs, N.J.: Prentice Hall. p. 231. ISBN 0-13-825001-4.*

A.41: Population Count. Programming Note.

6. ^ *Blaxell, David* (1978), "Record linkage by bit pattern matching", in *Hogben, David; Fife, Dennis W.*, [\*Computer Science and Statistics--Tenth Annual Symposium on the Interface\*](#), NBS Special Publication **503**, U.S. Department of Commerce / National Bureau of Standards, pp. 146–156.
7. ^ *Wegner, Peter* (1960), "A technique for counting ones in a binary computer", [\*Communications of the ACM\* \*\*3\*\* \(5\): 322](#), [doi:10.1145/367236.367286](#)
8. ^ "[GCC 3.4 Release Notes](#)" GNU Project
9. ^ "[LLVM 1.5 Release Notes](#)" LLVM Project.
10. ^ "[GHC 7.4.1 release notes](#)". GHC documentation.
11. ^ "[12.11. Bit Functions — MySQL 5.0 Reference Manual](#)".

## External links

- [Aggregate Magic Algorithms](#). Optimized population count and other algorithms explained with sample code.
- [HACKMEM item 169](#). Population count assembly code for the PDP/6-10.
- [Bit Twiddling Hacks](#) Several algorithms with code for counting bits set.
- [Necessary and Sufficient](#) - by Damien Wintour - Has code in C# for various Hamming Weight implementations.
- [Best algorithm to count the number of set bits in a 32-bit integer?](#) - Stackoverflow